

## NOTE 6635: Artificial Intelligence for Business Research

# Inference

Renyu (Philip) Zhang

1

## OpenAI API Prices

OpenAI Prices: <https://openai.com/api/pricing/>

Similar pricing scheme for Claude, Grok, DeepSeek, Qwen, etc.

<b>OpenAI o1</b> Frontier reasoning model that supports tools, Structured Outputs, and vision   200k context length  <b>Price</b> Input: \$16.00 / 1M tokens Cached input: \$7.50 / 1M tokens Output: \$60.00 / 1M tokens	<b>OpenAI o3-mini</b> Small cost-efficient reasoning model that's optimized for coding, math, and science, and supports tools and Structured Outputs   200k context length  <b>Price</b> Input: \$10 / 1M tokens Cached input: \$0.55 / 1M tokens Output: \$4.40 / 1M tokens
---	--

Save 50% on inputs and outputs with the [Batch API](#) and run tasks asynchronously over 24 hours.



@DrJimFan

<b>GPT-4.5</b> Largest GPT model designed for creative tasks and agentic planning, currently available in a research preview.   128k context length  <b>Price</b> Input: \$75.00 / 1M tokens Cached input: \$37.50 / 1M tokens Output: \$150.00 / 1M tokens	<b>GPT-4o</b> High-intelligence model for complex tasks   128k context length  <b>Price</b> Input: \$2.50 / 1M tokens Cached input: \$1.25 / 1M tokens Output: \$10.00 / 1M tokens	<b>GPT-4o mini</b> Affordable small model for fast, everyday tasks   128k context length  <b>Price</b> Input: \$0.150 / 1M tokens Cached input: \$0.075 / 1M tokens Output: \$0.600 / 1M tokens
---	--	---

### ChatGPT's User Experience: What is Behind the Decline in Intelligence?

ChatGPT | Bugs | chatgpt

**daixin0906** Jan 6 1 / 19 Jan 6

Since the beginning of this year, I have noticed some significant changes in the functionality and performance of ChatGPT, especially in terms of its intelligence and depth of reasoning. Once, whether as a work assistant or a daily conversational partner, ChatGPT left a deep impression on me. But now, with the use of the GPT-4o model and GPT-4o, I can't help but feel that their performance is far below the previous versions. This article will discuss this change from several perspectives.

<https://community.openai.com/t/chatgpts-user-experience-what-is-behind-the-decline-in-intelligence/1081511>

- **Key question:** How to make LLM inference more **efficient** and **cost-effective**?

2

2

# Agenda

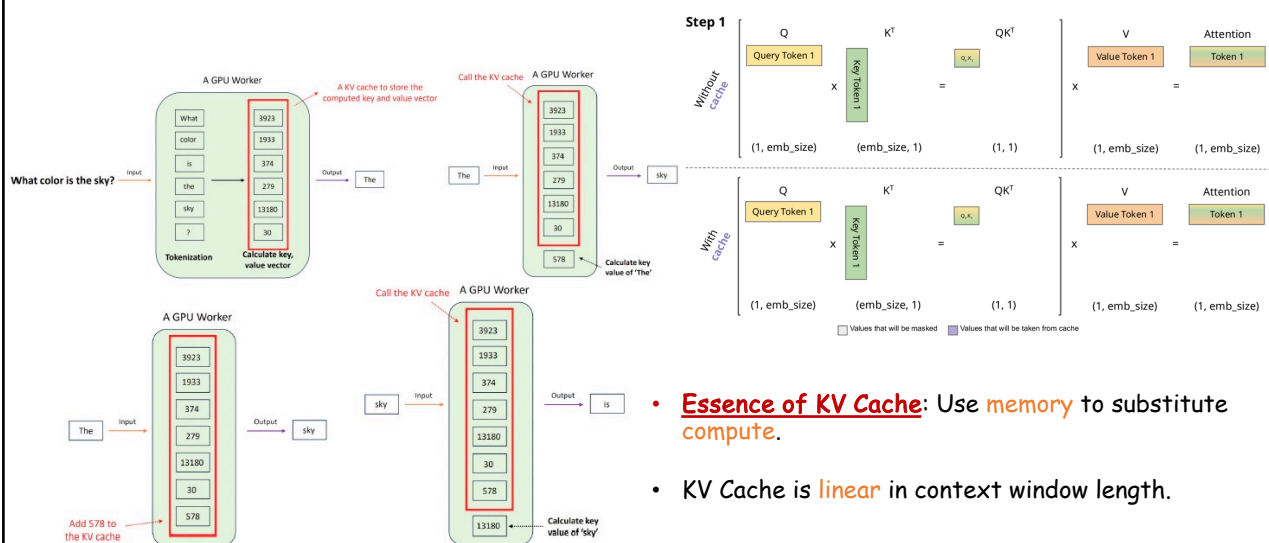
- KV-Cache
- Quantization
- DeepSeek Inference System
- OR for LLM Inference

3

3

## KV Cache

Hugging Face KV Cache Intro: <https://huggingface.co/blog/not-lain/kv-caching>  
<https://medium.com/@plienhar/llm-inference-series-2-the-two-phase-process-behind-llms-responses-1ff1ff021cd5>

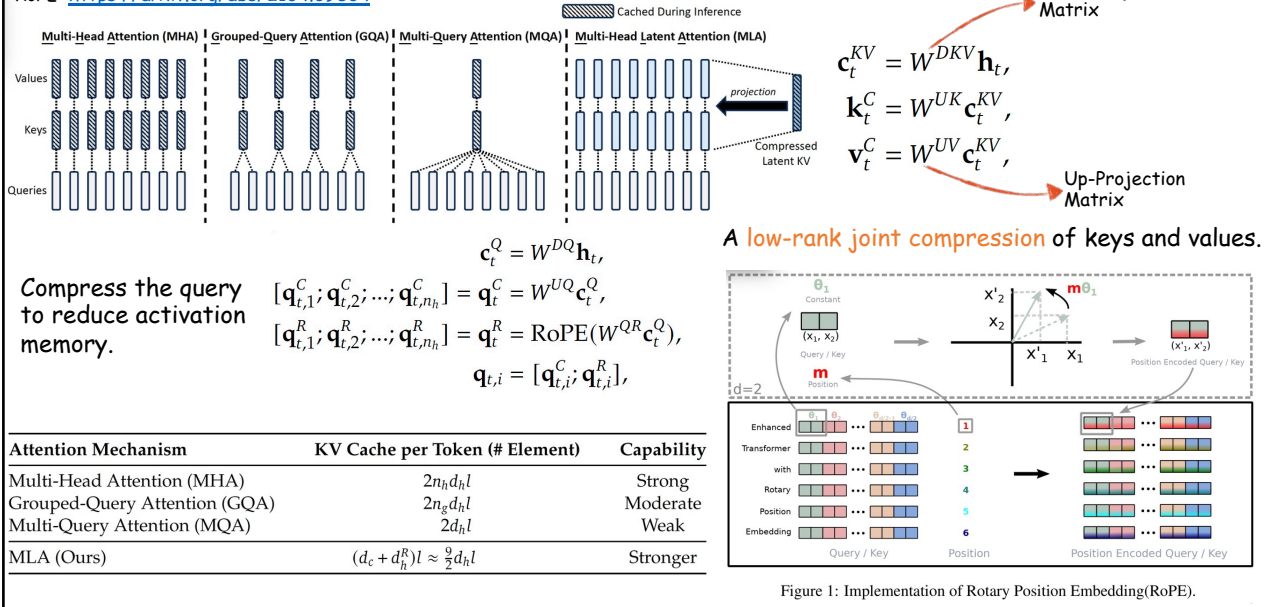


4

4

## Multi-Head Latent Attention

DeepSeek-V3: <https://arxiv.org/abs/2412.19437v1>; <https://www.bilibili.com/video/BV1HqFQezEMt>  
RoPE: <https://arxiv.org/abs/2104.09864>



5

## Agenda

- KV-Cache
- Quantization
- DeepSeek Inference System
- OR for LLM Inference

6

6

## Quantization

MIT Efficient DL Computing: <https://hanlab.mit.edu/courses/2024-fall-65940>

Quantization Fundamentals with Hugging Face: <https://learn.deeplearning.ai/courses/quantization-fundamentals/>


- **Quantization:** Mapping an input from a large (and continuous) set of values to a smaller (and discrete) set of values.
  - We do quantization to **save memory and energy** and **accelerate compute**, especially for LLM inference.



Data Type	torch.dtype
8-bit signed integer	torch.int8
8-bit unsigned integer	torch.uint8
16-bit signed integer	torch.int16
32-bit signed integer	torch.int32
64-bit signed integer	torch.int64

- For unsigned integer data types,  $[0, 2^n-1]$ .
- For signed integer data types,  $[-2^{n-1}, 2^{n-1}-1]$ .

## Two-Complement Representation



$$\begin{array}{ccccccc} 2^7 & + & 0 & + & 0 & + & 0 & + & 2^3 & + & 0 & + & 0 & + & 2^0 & = & 137 \\ 128 & & & & & & & & 8 & & & & & & 1 & & \end{array}$$



$-2^7 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0 = -49$

7

7


## Floating Number Representations

MIT Efficient DL Computing: <https://hanlab.mit.edu/courses/2024-fall-65940>

Quantization Fundamentals with Hugging Face: <https://learn.deeplearning.ai/courses/quantization-fundamentals/>

- Floating-point numbers:
  - Sign: +/-
  - Exponent: Range
  - Fraction/mantissa: Precision

Data Type	torch.dtype	torch.dtype alias
16-bit floating point	torch.float16	torch.half
16-bit brain floating point	torch.bfloat16	
32-bit floating point	torch.float32	torch.float
64-bit floating point	torch.float64	torch.double



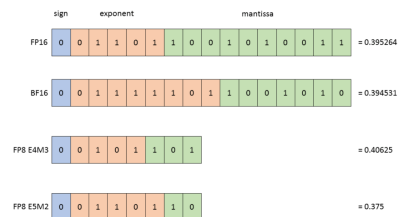

How to represent **0.265625**?

$$0.265625 = 1.0625 \times 2^{-2} = (1 + 0.0625) \times 2^{125-127}$$



Exponent =  $2^0 = 1$

Exponent  $\rightarrow 2^0 = 1$



8


8

## Linear Quantization

MIT Efficient DL Computing: <https://hanlab.mit.edu/courses/2024-fall-65940>

Quantization Fundamentals with Hugging Face: <https://learn.deeplearning.ai/courses/quantization-fundamentals/>

- Use a linear mapping to represent a number in high-precision type (FP32) in low-precision type (INT8).




9

## Linear Quantization

MIT Efficient DL Computing: <https://hanlab.mit.edu/courses/2024-fall-65940>

Quantization Fundamentals with Hugging Face: <https://learn.deeplearning.ai/courses/quantization-fundamentals/>

- How do we determine the scale  $s$  and zero point  $z$ ?




$$r_{\max} = S(q_{\max} - Z)$$

$$r_{\min} = S(q_{\min} - Z)$$


$$r_{\max} - r_{\min} = S(q_{\max} - q_{\min})$$

$$S = \frac{r_{\max} - r_{\min}}{q_{\max} - q_{\min}}$$

$$Z = \text{round}\left(q_{\min} - \frac{r_{\min}}{S}\right)$$



- Per-channel and per-group quantization.
- Quantization of weights and activations.
- Quantization-aware training.



10

## Quantization + LoRA: QLoRA

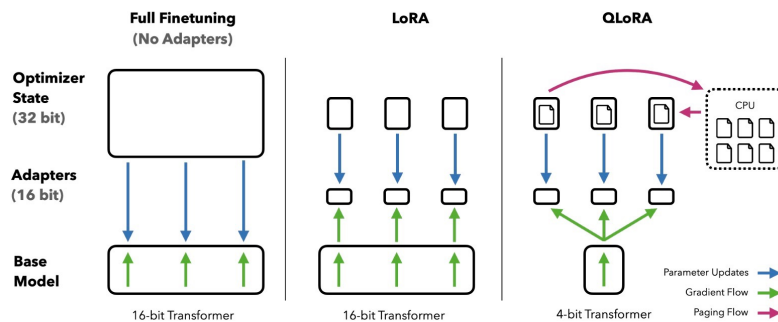
Stanford CS224N: <https://web.stanford.edu/class/cs224n/>  
 QLoRA Paper: <https://arxiv.org/abs/2305.14314>

**Qlora: Efficient finetuning of quantized llms**

T. Dettmers, A. Pagnoni, A. Holtzman... - Advances in neural ..., 2023 - proceedings.neurips.cc

... We present **QLORA**, an efficient finetuning approach that reduces ... **QLORA** backpropagates gradients through a frozen, 4-bit ... **QLORA** introduces a number of innovations to save memory ...


☆ Save ⓘ Cite Cited by 2449 Related articles All 9 versions ⓘ



**Figure 1:** Different finetuning methods and their memory requirements. QLoRA improves over LoRA by quantizing the transformer model to 4-bit precision and using paged optimizers to handle memory spikes.

To further save memory, adopt **double-quantization (DQ)**.

- QLoRA improves over LoRA by **quantizing** the transformer to **4-bit precision** and using **paged optimizer** to handle memory.
- 4-bit NormalFloat (NF4)
  - Data type suitable for **normally distributed** weights.



11

11

## Agenda

- KV-Cache
- Quantization
- DeepSeek Inference System
- OR for LLM Inference


12

12

## DeepSeek-V3/R1 Inference System

GitHub: [https://github.com/deepseek-ai/open-infra-index/blob/main/202502OpenSourceWeek/day\\_6\\_one\\_more\\_thing\\_deepseekV3R1\\_inference\\_system\\_overview.md](https://github.com/deepseek-ai/open-infra-index/blob/main/202502OpenSourceWeek/day_6_one_more_thing_deepseekV3R1_inference_system_overview.md)

- Extremely optimized for high throughput and low latency: cross-node Expert Parallelism (EP).
  - Leverage EP to scale batch size.
  - Hide communication latency behind computation.
  - Perform load balancing.
- Served with 278 8-H800 GPU nodes; average occupancy 226.75 nodes; each with throughput ~73.7k tokens/s for input during prefilling and ~14.8k tokens/s for output during decoding.
- Daily input tokens: 608B (342B hit the KV cache)
- Daily output tokens: 168B; 20-22 tokens/s; average KV-cache length per output: 4,989 tokens.



Daily cost = \$87,072; Daily Revenue under the R1-API pricing = \$562,027, i.e., **545% profit margin**

13

13


## Agenda

- KV-Cache
- Quantization
- DeepSeek Inference System
- OR for LLM Inference

14

14

# OR for LLM Inference



## Fundamental Modeling for LLM Inference with Exploding KV Cache Demands


Patrick Jaillet  
Department of Electrical Engineering and Computer Science, Massachusetts Institute of Technology, jaillet@mit.edu

Jiahuo Jiang  
HKUST, jsjiang@ust.hk


Chara Podimata  
Sloan School of Management, Massachusetts Institute of Technology, podimata@mit.edu

Zijie Zhou\*  
Operations Research Center, Massachusetts Institute of Technology, zhou98@mit.edu

- Given the KV-cache memory limit, how to **batch** different prompts and output tokens to **minimize total end-to-end latency**.
- Proposed scheduling algorithm (with **provable constant regret**): Prioritize the prompt with the smallest **predicted** number of output tokens, subject to the KV-cache limit constraint.
  - Benchmark: alpha-protection first-come-first-serve, beta-clearing when reaching KV-cache limit.



High-Demand



Low-Demand

15